

# Computational Optimal Transport - In a Nutshell

Main Reference: Computational Optimal Transport

Key Concepts:

Optimal Transport (OT)

Sinkhorn's Algorithm

Wasserstein Distance

Sliced Wasserstein Distance

Table of Contents:

- ⌚ Problem Overview: What Is Optimal Transport?
- ✖ Discrete Optimal Transport
- 📦 Sinkhorn Algorithm
- ⚡ Sliced Wasserstein Distance: A Fast High-Dimensional Alternative

## Problem Overview: What Is Optimal Transport?

Optimal Transport (OT) provides a geometric and principled framework for comparing probability distributions.

At its core, the OT problem asks:

How can we transform one distribution into another while incurring the least possible cost?

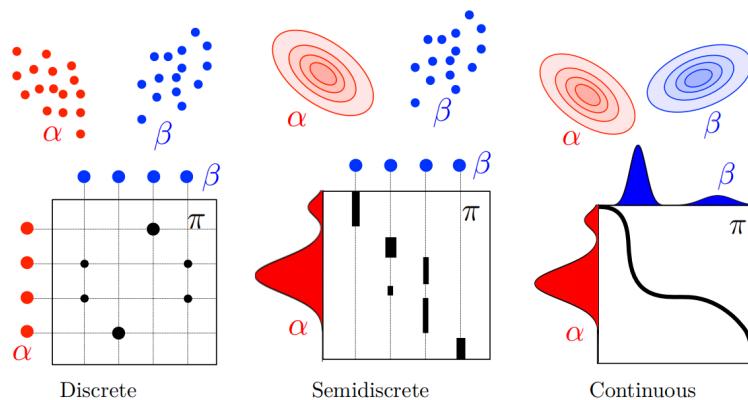
Formally, given two probability measures  $\mu$  and  $\nu$ , OT seeks to determine how much mass should be moved from each location in  $\mu$  to each location in  $\nu$  in the most efficient way.

In the Kantorovich formulation, we search for a transport plan  $\pi(x, y)$ , which specifies how much mass is moved from  $x$  to  $y$ . The goal is to minimize the total cost:

$$\min_{\pi \in \Pi(\mu, \nu)} \int c(x, y) d\pi(x, y),$$

where  $c(x, y)$  is the cost of transporting a unit of mass from  $x$  to  $y$ , and  $\Pi(\mu, \nu)$  denotes the set of all couplings whose marginals are  $\mu$  and  $\nu$ .

Although OT can be defined for continuous, discrete, or mixed (semi-discrete) measures, in this article we focus on the discrete setting, which is both intuitive and central to computational OT.



# Discrete Optimal Transport

We begin with the discrete setting, which is both intuitive and widely used in practical applications.

Assume that the two distributions are represented as

$$\mu = \sum_i \mathbf{a}_i \delta_{x_i}, \quad \nu = \sum_j \mathbf{b}_j \delta_{y_j},$$

where  $\mathbf{a}_i$  and  $\mathbf{b}_j$  are the associated masses, and

$$\mathbf{C}_{ij} = c(x_i, y_j)$$

defines the pairwise transportation cost.

We seek a transport matrix  $\mathbf{P} \in \mathbb{R}^{n \times m}$  that satisfies the mass-preserving constraints:

$$\mathbf{P}_{ij} \geq 0, \quad \sum_j \mathbf{P}_{ij} = \mathbf{a}_i, \quad \sum_i \mathbf{P}_{ij} = \mathbf{b}_j.$$

Under these constraints, the discrete OT problem becomes the following linear program:

$$\min_{\mathbf{P}} \sum_{i,j} \mathbf{C}_{ij} \mathbf{P}_{ij}.$$

Each entry  $\mathbf{P}_{ij}$  specifies how much mass is transported from location  $x_i$  to location  $y_j$ . The total transport cost is simply the weighted sum of all such transport assignments.

 **Wasserstein Distance.** When the ground cost captures a meaningful notion of distance between points, most commonly

$$c(x, y) = \|x - y\|^p,$$

the optimal transport objective directly inherits this geometric structure. In this setting, the minimal transport cost defines the  **$p$ -Wasserstein distance**:

$$W_p(\mu, \nu) := \left( \min_{\mathbf{P} \in \Pi(\mu, \nu)} \sum_{i,j} \|x_i - y_j\|^p \mathbf{P}_{ij} \right)^{1/p}.$$

In one dimension, Wasserstein distances become even simpler: if the samples of  $\mu$  and  $\nu$  are sorted, the formula reduces to a point-wise difference. Formally, given two empirical 1D measures

$$\alpha = \frac{1}{n} \sum_{i=1}^n \delta_{x_i}, \quad \beta = \frac{1}{n} \sum_{i=1}^n \delta_{y_i},$$

and assuming the points are sorted

$$x_1 \leq \dots \leq x_n, \quad y_1 \leq \dots \leq y_n,$$

the  $p$ -Wasserstein distance reduces to:

$$W_p(\alpha, \beta)^p = \frac{1}{n} \sum_{i=1}^n |x_i - y_i|^p.$$

Hence, In one dimension, the Wasserstein distance has a closed-form expression—just sort the points.

# Sinkhorn Algorithm

Classical optimal transport solved via linear programming is computationally expensive, typically requiring  $O(n^3)$  time. This quickly becomes impractical for large datasets such as images, point clouds, or minibatches in machine learning. To address this limitation, the Sinkhorn algorithm introduces an entropy-regularized formulation:

$$\min_{\mathbf{P}} \sum_{i,j} \mathbf{C}_{ij} \mathbf{P}_{ij} + \varepsilon \sum_{i,j} \mathbf{P}_{ij} (\log \mathbf{P}_{ij} - 1).$$

Here the first term represents the transport cost, the second term encourages  $\mathbf{P}$  to stay diffuse (high entropy), and  $\varepsilon > 0$  controls the smoothness-accuracy trade-off.

 **Factorized Solution.** To enforce the marginal constraints, we introduce Lagrange multipliers:

- a vector  $\boldsymbol{\alpha} \in \mathbb{R}^n$  for the row constraints ( $\mathbf{P}\mathbf{1} = \mathbf{a}$ ),
- a vector  $\boldsymbol{\beta} \in \mathbb{R}^m$  for the column constraints ( $\mathbf{P}^\top \mathbf{1} = \mathbf{b}$ ).

The Lagrangian of the regularized problem can be written as

$$\mathcal{L}(\mathbf{P}, \boldsymbol{\alpha}, \boldsymbol{\beta}) = \sum_{i,j} \mathbf{C}_{ij} \mathbf{P}_{ij} + \varepsilon \sum_{i,j} \mathbf{P}_{ij} (\log \mathbf{P}_{ij} - 1) + \sum_i \alpha_i \left( a_i - \sum_j \mathbf{P}_{ij} \right) + \sum_j \beta_j \left( b_j - \sum_i \mathbf{P}_{ij} \right).$$

Taking the derivative of  $\mathcal{L}$  with respect to  $\mathbf{P}_{ij}$  and setting it to zero (stationarity) yields

$$\varepsilon \log \mathbf{P}_{ij} = \alpha_i + \beta_j - \mathbf{C}_{ij}.$$

Exponentiating both sides, we obtain

$$\mathbf{P}_{ij} = \exp(\alpha_i/\varepsilon) \exp(-\mathbf{C}_{ij}/\varepsilon) \exp(\beta_j/\varepsilon).$$

Now define the scaling vectors and kernel matrix

$$\mathbf{u}_i := e^{\alpha_i/\varepsilon}, \quad \mathbf{v}_j := e^{\beta_j/\varepsilon}, \quad \mathbf{K}_{ij} := e^{-\mathbf{C}_{ij}/\varepsilon},$$

so that the optimal transport plan takes the factorized form

$$\mathbf{P} = \text{diag}(\mathbf{u}) \mathbf{K} \text{diag}(\mathbf{v}).$$

 **Sinkhorn Iterations.** The scaling vectors  $\mathbf{u}$  and  $\mathbf{v}$  are computed via alternating normalization, also known as the Sinkhorn updates:

```
# Sinkhorn iterations
u = 1
while not converged:
    v = b / (K.T @ u)      # column normalization
    u = a / (K @ v)        # row normalization
```

Each step enforces one set of marginal constraints:

- $\mathbf{v} \leftarrow \mathbf{b} \oslash (\mathbf{K}^\top \mathbf{u})$  ensures  $\mathbf{P}^\top \mathbf{1} = \mathbf{b}$
- $\mathbf{u} \leftarrow \mathbf{a} \oslash (\mathbf{K} \mathbf{v})$  ensures  $\mathbf{P} \mathbf{1} = \mathbf{a}$

where  $\oslash$  denotes elementwise division. These iterations converge under mild conditions and are highly efficient with cost  $O(n^2)$ .

# Sliced Wasserstein Distance: A Fast High-Dimensional Alternative

Computing Wasserstein distances in high dimensions is notoriously difficult. For empirical measures on  $X = \mathbb{R}^d$  (with bounded support), the statistical convergence rate is

$$\mathbb{E} |W_p(\hat{\alpha}_n, \hat{\beta}_n) - W_p(\alpha, \beta)| = O(n^{-1/d}),$$

which rapidly worsens as the dimension  $d$  increases. This *curse of dimensionality* makes both estimation and computation of OT challenging in high-dimensional settings.

 **Sliced Wasserstein Distance (SWD).** The SWD provides an elegant workaround: instead of operating directly in  $d$ -dimensions, it repeatedly **projects the distributions onto 1D lines**, computes efficient **1D Wasserstein distances**, and aggregates the results. This preserves much of the geometry of OT while being significantly faster and statistically more robust.

In 1D, optimal transport is remarkably simple:

$$W_p(\alpha, \beta)^p = \frac{1}{n} \sum_{i=1}^n |x_{(i)} - y_{(i)}|^p,$$

where  $x_{(i)}$  and  $y_{(i)}$  denote sorted samples. This closed form makes 1D OT extremely efficient.

Given two distributions  $\mu, \nu \subset \mathbb{R}^d$ , SWD proceeds as follows:

## 1. Sample random directions

$$\theta \sim \text{Uniform}(S^{d-1})$$

## 2. Project samples onto $\theta$

$$x \mapsto \langle x, \theta \rangle$$

## 3. Compute 1D Wasserstein distance

Sort the projected samples and apply the 1D formula.

## 4. Repeat and average across directions

A Monte Carlo integral approximation.

The resulting sliced Wasserstein distance is

$$SW_p(\mu, \nu) = \left( \int_{S^{d-1}} W_p(P_\theta \mu, P_\theta \nu)^p d\theta \right)^{1/p},$$

and in practice is approximated via

$$SW_p(\mu, \nu) \approx \left( \frac{1}{L} \sum_{\ell=1}^L W_p(P_{\theta_\ell} \mu, P_{\theta_\ell} \nu)^p \right)^{1/p}.$$

Although SWD does not completely eliminate high-dimensional effects, its statistical behavior is dramatically better: empirical rates are often close to  $O(n^{-1/2})$ , independent of the ambient dimension. This makes SWD a highly attractive OT surrogate for applications in imaging, generative modeling, and high-dimensional statistics.